

# CS553: Cloud Computing

## Evaluation Document


### CPU Benchmarking:

---

- System Configuration
  - Amazon EC2 – t2-micro
  - Operating System: Linux Ubuntu 14\*
  - RAM: 1 GB
  - 1 Core

Number Of Threads	GFLOPS	GIOPS
1	8.646125	8.669155
2	8.639707	8.130392
3	8.657069	8.660769

- Above Table shows CPU speed in terms of GIGA FLOPS (Giga Floating Point Operation per second) and GIGA IOPS (Giga Integer Operations per second respectively)
- We have taken 3 different threads (1,2, and 4).
- We have received almost same readings for GFLOPS and GIOPS irrespective of number of threads.
- But we have observed as we keep increasing number of threads to higher number performance start decreasing as threads share CPU and they need to switch among themselves more frequently than threads count with optimal number.




- **CPU Theoretical Performance –**

$$\begin{aligned}
 &= \text{CPU Frequency} * \text{Number of Cores} * \text{Threads per core} * \text{IPC} \\
 &= 2.69 * 1 * 1 * 16 \\
 &= 43.04
 \end{aligned}$$


- We have calculated CPU GFLOPS 8.6 and Theoretical Value for CPU is 43.04, so we have achieved 19.98% of Theoretical performance.

- GFLOPS for 10 min (sample plotted for every second)



- We keep calculating GFLOPS over the period of 10 min (600 seconds) and keep the track of GFLOPS per second, Performance is fluctuating and giving the values between 10 to 30 GFLOPS on an average

- GILOPS for 10 min (sample plotted for every second)



- We keep calculating GLOPS over the period of 10 min (600 seconds) and keep the track of GLOPS per second, Performance is fluctuating and giving the values between 10 to 35 GFLOPS on an average

### Conclusion-

- If we keep the execution running for 600 seconds and more, we found slight increase in the average number of GFLOPS and GLOPS.
- **Extra Credit:**
  - Here we have compared our results with Linpack benchmark. Below is screen shot for Linpack systems outputs

```

the correct number of CPUs/threads, problem input files, etc..
Fri Feb 12 00:53:14 UTC 2016
Intel(R) Optimized LINPACK Benchmark data

Current date/time: Fri Feb 12 00:53:14 2016

CPU frequency:      2.835 GHz
Number of CPUs: 1
Number of cores: 1
Number of threads: 1

Parameters are set to:

Number of tests: 15
Number of equations to solve (problem size) : 1000  2000  5000  10000 15000 1800
0 20000 22000 25000 26000 27000 30000 35000 40000 45000
Leading dimension of array      : 1000  2000  5008  10000 15000 1800
8 20016 22008 25000 26000 27000 30000 35000 40000 45000
Number of trials to run
  2    2    2    2    1    1    1    1    1
Data alignment value (in Kbytes) : 4    4    4    4    4    1    1    1    1
  4    4    4    4    4    1    1    1    1

Maximum memory requested that can be used=800204096, at the size=10000

===== Timing linear equation system solver =====

Size  LDA  Align. Time(s)  GFlops  Residual  Residual(norm) Check
1000  1000  4      0.039  17.1631  9.900691e-13  3.376390e-02  pass
1000  1000  4      0.037  18.0275  9.900691e-13  3.376390e-02  pass
1000  1000  4      0.037  18.1542  9.900691e-13  3.376390e-02  pass
1000  1000  4      0.037  17.9963  9.900691e-13  3.376390e-02  pass
2000  2000  4      0.276  19.3540  4.053480e-12  3.526031e-02  pass
2000  2000  4      0.275  19.3988  4.053480e-12  3.526031e-02  pass
5000  5008  4      4.009  20.7975  2.336047e-11  3.257429e-02  pass
5000  5008  4      4.004  20.8267  2.336047e-11  3.257429e-02  pass
10000 10000 4     31.436  21.2132  1.124127e-10  3.963786e-02  pass
10000 10000 4     30.908  21.5757  1.124127e-10  3.963786e-02  pass

Performance Summary (GFlops)

Size  LDA  Align.  Average  Maximal
1000  1000  4       17.8353  18.1542
2000  2000  4       19.3764  19.3988
5000  5008  4       20.8121  20.8267
10000 10000 4       21.3944  21.5757


Residual checks PASSED

End of tests

Done: Fri Feb 12 00:54:39 UTC 2016
ubuntu@ip-172-31-56-58:/usr/share/linpack$

```

- Analysis-
  - Below is graphical representation of results of my system with Linpack
  - My system has average GLOPS value to 8.6, Whereas Linpack has 20.5
  - So here we have achieved almost 42% of ideal value compared with Linpack.





## Memory Benchmarking:

- Write + Read Operations (Throughput)

	Sequential	
Block Size	Thread 1	Thread 2
1 Byte	201.207243	204.081633
1 KB	4266.66667	4876.19048
1 MB	6990.50667	7767.22963

	Random	
Block Size	Thread 1	Thread 2
1 Byte	8.719916	8.245723
1 KB	2275.55556	2178.7234
1 MB	8065.96923	7489.82857





- Analysis-

- From above tables and graph we can conclude that throughput for Read+Write Operations goes on increasing with increase in Size of block we are reading and writing by 1 Thread or 2 Threads
- This is because as we increase block size, CPU write or Read more data at a time cause increase in throughput. It is not dependent on number of thread performing operation.

- Write + Read Operations (Throughput)

	Thread 1	
Block Size	Sequential	Random
1 Byte	201.207243	8.719916
1 KB	4266.66667	2275.55556
1 MB	6990.50667	8065.96923



	Thread 2	
Block Size	Sequential	Random
1 Byte	204.081633	8.245723
1 KB	4876.19048	2178.7234
1 MB	7767.22963	7489.82857



- Analysis –
  - From above tables and graph we can conclude that throughput for Read+Write Operations goes on increasing with increase in Size of block we are reading and writing it Sequentially OR Randomly.
  - This is because as we increase block size, CPU write or Read more data at a time cause increase in throughput. It is not dependent on how we are performing operations like Sequential or Random, even there is slight difference in their operations per seconds.
- Write + Read Operations (Latency)

	Thread 1	
Block Size	Sequential	Random
1 Byte	4.97	114.68
1 KB	0.234375	0.439453
1 MB	0.143051	0.123978

	Thread 2	
Block Size	Sequential	Random
1 Byte	4.9	121.275
1 KB	0.205078	0.458984
1 MB	0.128746	0.133514



- Analysis –
  - Tables and graphs above showing Latency for Sequential and Random operations for 1 Thread and 2 Threads.
  - Here we can conclude that Latency is high when we are reading or Writing Randomly with block Size of 1 Byte, in both the cases i.e. 1 thread or 2 thread.

- **Memory Theoretical Performance –**

= Clock Frequency \* Number of data transfer per clock \* Memory bus interface width \* number of interfaces

= 1600MHz \* 2 \* 64 bits \* 2

= 409600 Byte

= 51200 MB

= 51200 MHz

= 51.2 GHz

- **Extra Credit**

- Here we have compared our results with iotzone benchmark. Below is screen shot for iotzone systems outputs.

```
Run began: Fri Feb 12 01:00:04 2016

Auto Mode
Command line used: ./iotzone -a
Output is in Kbytes/sec
Time Resolution = 0.000001 seconds.
Processor cache size set to 1024 Kbytes.
Processor cache line size set to 32 bytes.
File stride size set to 17 * record size.
```


	KB	reclen	write	rewrite	read	reread	random read	random write	bkwd read	record rewrite	stride read	fwrite	frewrite	fread	freread
64	4	855419	2006158	10821524	10402178	7100397	3203069	4988978	4018152	7100397	3203069	3363612	5860307	10821524	
64	8	1421755	2203800	7100397	15972885	12902017	2662899	4274062	4274062	5389653	2561267	2662899	4897948	12310336	
64	16	1049372	2298136	3203069	15972885	12310336	2662899	4274062	3363612	7940539	3057153	3165299	5389653	7940539	
64	32	1330167	2467108	9318832	15972885	10821524	2801873	4564786	3057153	5735102	3203069	3588436	4564786	12310336	
64	64	1452528	2358717	15972885	15972885	10402178	2662899	4274062	2772930	4018152	2689580	4274062	7940539	12902017	
128	4	1172704	2409592	11720614	11720614	8548124	3657016	7082197	4717434	8414153	3657016	3657016	8036304	9795896	
128	8	1598754	3657016	11720614	14200794	11720614	3982553	7917784	5122535	8548124	3657016	4012317	9129573	12542043	
128	16	1727352	4407601	14200794	14200794	12542043	4135958	8414153	4596273	8036304	3867787	4557257	9795896	10779307	
128	32	1827299	4557257	12842051	5784891	11470204	4717434	8036304	5122535	10779307	4267461	4557257	9129573	12842051	
128	64	1852520	4407601	14200794	15881078	12842051	4717434	9129573	4717434	7582312	3218540	6727225	9129573	12842051	
128	128	1663139	4596273	5784891	15881078	12842051	4934216	8548124	4759253	5325799	3468030	4557257	8548124	12542043	
256	4	1347557	3705040	8534922	10651598	8534922	3159869	5810112	4652146	5455847	3605511	3511189	8208677	9434868	
256	8	1407621	4264168	12228612	13454450	10758322	3935921	9518507	4197489	9868433	3421677	1938842	3654598	13454450	
256	16	1675612	2413957	7314033	13454450	12228612	3087188	6107548	5569035	11569783	2582316	2812272	4070199	12228612	
256	32	1354356	2812272	5455847	12812277	12228612	2975955	5569035	5455847	9518507	3159869	3087188	6398720	12812277	
256	64	1588832	2509882	7314033	15164624	12661199	3236056	5687020	4929815	8815202	2842047	3009318	3823789	13454450	
256	128	1610276	3009318	8271916	14164395	12228612	3052087	4477549	3009318	4281170	2754556	3756894	4840911	12812277	
256	256	1673001	2911402	10758322	7314033	7571923	3326279	4929815	2486631	3705040	2665657	2000242	2984226	12812277	
512	4	1160923	2073249	5886650	8844453	8528336	2572435	5177083	4030519	6571132	2032051	2275345	4742616	10235583	
512	8	1418599	2089386	6571132	14240069	10485468	2600470	5951911	5551840	9859630	3029721	2651850	3963567	12499490	
512	16	1514652	2483197	6821616	13783088	11374040	3372274	5227492	6228097	9304292	2350044	2317080	6821616	12797441	
512	32	1723770	2786024	7115451	14240069	11877300	2681653	6246213	6319739	10044089	2906696	2216629	4784885	13438092	
512	64	1610042	2317080	6931711	13109944	11877300	3304808	5384786	5278892	11374040	2826358	2995907	6472112	13872122	
512	128	1689859	2497638	7647578	1352271	1113109944	3580298	4871723	3970896	8992598	3220553	3659616	6319739	13872122	
512	256	1815584	2235086	7115451	12215097	9107004	3029721	5509113	3346002	4494470	2497638	4340054	5566231	9468384	
512	512	1595686	2497638	11620224	12499490	5214798	3008498	6086873	2926501	5127637	3393590	3068685	6319739	8528336	
1024	4	1182598	1858644	5720477	11645609	7420395	2403712	5303701	5564829	8457895	2760606	2197132	4720752	9485232	
1024	8	1607514	2716948	6692005	12639479	8457895	2598591	5917516	6863100	10329265	2625597	1954207	6059442	11018226	
1024	16	1478064	2290886	5821271	11520658	11645609	2960401	6093831	6402699	9485232	2716948	2173780	5853003	7220790	
1024	32	1648862	2348509	6819511	10662628	12037272	2701567	5821271	6787181	9677584	2918161	2659742	7008693	13863422	
1024	64	1651398	2960401	6744549	14230902	12828238	3529706	6918376	5720477	13472053	3066069	3191372	6519323	14668318	
1024	128	1503415	2751762	7208671	13305116	10878686	2966535	5885083	4762630	10354166	3425544	2497356	6830356	12071103	
1024	256	1625154	2579861	7869040	12983353	11903823	2944166	6650556	5390231	8914314	3655895	3458646	6819511	8895850	
1024	512	1865101	2297012	9402175	12983353	8840915	3618930	5958564	4594502	8822754	3149251	4826859	6874084	11249091	
1024	1024	1723636	2491561	10796646	13142265	12313351	3191372	6819511	3281592	6519323	3314514	2640123	6787181	12983353	
2048	4	1354493	2321725	5626082	11502234	9101380	2283460	5293261	5785224	8292998	2438378	1905146	5159722	10607402	
2048	8	1445434	2120026	6807652	10240672	11202231	3102397	6202989	7393606	11365268	3169944	2572756	6696208	12882215	
2048	16	1605335	2003637	6017202	13201113	12187663	2387638	6114678	6235466	11000610	3042063	2767540	6308502	12488807	



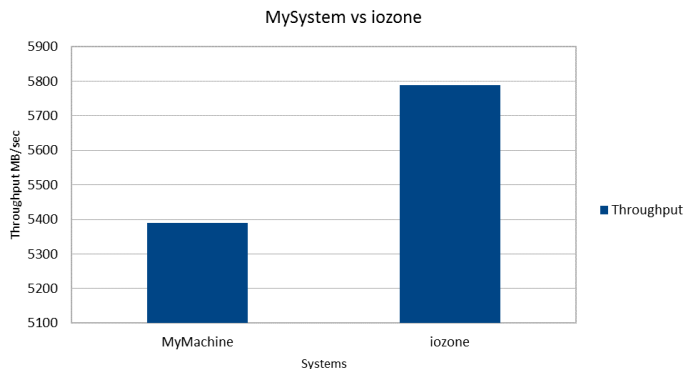
16384	16384	1723038	2491301	10736840	15422281	12311388	3493972	8083944	3281932	8713325	3447428	1008123	8787181	12782333
2048	4	1354493	2321725	5626082	11502234	9101380	2283460	5293261	5785224	8202998	2438378	1005146	5159722	10687402
2048	8	1445434	2120026	6807652	10240672	11202231	3102397	6202989	7393606	11365268	3169944	2572756	6696208	12882215
2048	16	1695335	3002627	6917293	1330111	312187663	3287628	6114678	6225466	11900619	2942962	2767549	6298503	12488897
2048	32	1563282	2476337	7311791	141326981	2959958	2967361	5611381	7496851	12729493	2813783	2639154	6382745	13038646
2048	64	1523901	2270783	6873016	1330111	313639023	3538147	6207471	6945257	11564174	3180508	2959183	6401772	12635867
2048	128	1713598	3089009	6608629	13748169	11440955	2774700	5738844	6649556	12488897	3250308	2892423	6917293	13489098
2048	256	1641240	3047366	5674397	10834854	11502234	3025897	5832361	6041567	9944290	3038742	2457914	6690992	11835033
2048	512	1806199	2933915	7698414	12187663	11851361	3199462	6543188	4994712	9578367	3730203	3271351	6261772	11980619
2048	1024	1731211	2709070	8749118	13038646	11564174	2812862	6543188	4385291	8987113	3319389	5071381	6945257	12729493
2048	2048	1958129	3180508	12882215	13138359	11770166	3292669	6416117	3671217	6167360	3549844	3006831	6266348	11502234
4096	4	1354045	1903416	6461498	10867793	8603675	4338563	5572903	5284803	8062616	3357537	2652931	6122971	9868918
4096	8	1882352	2992067	6758923	12880381	1252164	5263753	6522830	6682678	10136778	4485833	2505520	6041154	10118866
4096	16	1608179	2663212	7418527	11641010	10985935	3470816	6178018	7277121	11447097	3571100	2740095	6861502	11311433
4096	32	2131020	3116922	7625987	12967840	13258059	3800178	6847827	7135072	11378861	4095510	2940849	6169123	10698597
4096	64	1963457	3133979	7460407	12680688	13518879	4231697	7209932	7935986	11135471	3750403	2910951	6839648	13176709
4096	128	2058256	3199942	7656576	12569357	13340420	5828019	7249484	7261741	11833449	3935097	2940849	6169123	10698597
4096	256	1974968	2896717	7173805	11042425	12765490	4008550	7123239	6916752	13207098	3501232	2970386	6804432	11348794
4096	512	1955189	3286877	7249484	11508442	11907264	3314142	6500612	5859825	10190895	2730514	2940849	6169123	10698597
4096	1024	1923879	2878275	7892238	12523544	12083134	3045641	6638778	5145513	10089154	3734100	2977031	5019235	11164417
4096	2048	1703202	2483786	9163546	13476460	12919082	3540193	4954103	4887857	7892238	3447830	5050220	6714018	12007129
4096	4096	1718707	3009890	10215133	12117223	13086376	3667146	5312585	3374022	5196878	3212508	3185110	6575257	10779149
8192	4	1424958	2489149	6272536	8819674	8532752	2564005	4757117	5672914	8551866	2658022	2282461	5044886	9536939
8192	8	1485132	2226697	5285458	11554130	11735662	3119697	5178710	7154097	11084481	3093575	2612151	7013901	8111736
8192	16	1654311	2716226	7142200	11240406	12409588	4213951	6854392	6351381	10981742	3676880	2649414	7453619	11084481
8192	32	1863952	2722467	7218727	11457807	12468128	4162896	7543623	8047143	12778810	3777121	2675407	7321790	10530733
8192	64	1855497	2911150	7802283	10805609	12642455	4471708	7346838	7500806	10115308	3503662	2756540	6977775	10291026
8192	128	1781607	2718805	7155586	11102389	12199550	3494753	5224380	8102172	11652086	3304518	2398115	7068733	10569606
8192	256	1746382	3070628	6506533	11457807	12642455	3491912	6425016	7167528	12302946	3102794	2318032	6786699	10393756
8192	512	1792575	2792837	5919188	10908526	11804201	3422010	6336155	6186701	11393220	3143385	2888633	6559949	10315743
8192	1024	1612617	2544825	6670750	10569606	10738070	3691893	7295362	6215800	10925870	3410461	2926772	6770651	9905351
8192	2048	1730463	2788304	8242105	11620560	12860118	3479888	6687628	5378957	10625170	3551098	3665505	7944805	10501766
8192	4096	1908258	3280853	10422130	12942468	13065504	5984134	7578564	6262248	9383277	4158362	4364357	7772279	10751510
8192	8192	2006313	2957252	9002223	10315743	8577485	4625203	7224798	5264403	7367317	3978753	3364048	6224809	8725605
16384	4	1414860	2507478	6384903	7790612	7136122	2853347	5883111	5809993	8079126	2841079	2352294	6023899	6962596
16384	8	1531940	2646435	7192137	8881003	9204562	2811784	6537983	6551072	9372790	3016816	2547259	7147999	9153070
16384	16	1739477	2866320	7062054	9214436	10958635	3256722	6942900	7173368	9953854	2999042	2628116	7052632	8972610
16384	32	1706518	2866320	7356911	9841241	11260308	3240597	7010899	7703281	12128809	3095359	2803068	7595993	10063176
16384	64	1807988	3174578	7508843	10402891	112449618	3802158	6993063	7775626	9991482	3196878	2869791	7379822	9740803
16384	128	1790977	2988088	7403675	10200602	11949538	3679205	5893706	7702417	11314072	3518156	2732514	7201936	9654587
16384	256	1764761	2926249	6614125	8475718	9372790	4010753	6901759	6644181	10694297	3284743	2798274	5929815	8191806
16384	512	1746729	2673617	7245981	8687884	9438443	3019998	5917050	6939394	10618284	2889581	2679351	6716917	9834199
16384	1024	1833066	3119103	7278213	9592593	11643805	4487419	7656078	6806735	12109573	3385728	2912112	6787239	9097331
16384	2048	1815631	2823801	7585093	8890194	10436067	3402660	7028827	5748753	11488082	3229024	3205079	7016626	7865515
16384	4096	1763810	3041383	7950141	9966847	10577425	3946037	7353761	6097662	10182464	3715207	3425388	7349829	8991393
16384	8192	1762091	2895303	7854726	8769931	8445510	3612858	6566722	5287006	8245869	4179557	4176255	6480642	9157949
16384	16384	1701701	2915325	6592551	6854943	7080244	3232822	4787877	3209419	5069376	3102906	2924630	4817414	6363028
32768	64	1786991	3148672	7100291	8173200	8999522	3244024	7459047	8180010	9029083	3491949	2992565	7546695	8064811
32768	128	1749036	3177058	7340332	8092829	8533503	3400960	7403597	8019635	8573962	3196938	3072167	7391652	8312602

- Analysis –
  - Below is graphical representation of results of my system with iotzone
  - So here we have achieved almost below % of ideal value compared with iotzone
  - Sequential Write Operation = 71.80%
  - Sequential Read Operation = 91.11%
  - Random Write Operation = 40.45 %
  - Random Read Operation = 118.36%

- Sequential Operations




Sequential Write Operation




Sequential Read Operation



- Random Operations



Random Write Operation





Random Read Operation

## Disk Benchmarking:

- Write Operations (Throughput)

	Sequential	
Block Size	Thread 1	Thread 2
1 Byte	32.938076	32.615786
1 KB	1163.63636	1125.27473
1 MB	1978.44528	1807.88966

	Random	
Block Size	Thread 1	Thread 2
1 Byte	0.846762	0.845802
1 KB	572.067039	567.313019
1 MB	1978.44528	1823.61044





- Analysis –
  - Tables and graphs above are for Throughput of Write Operations for 1 thread and 2 threads.
  - From above graph we can conclude that throughput goes on increasing with increase in block size from 1 byte to 1 MB.
  - It increases in same fashion for both the operation i.e. with 1 thread and 2 thread.

- Read Operations (Throughput)

	Sequential	
Block Size	Thread 1	Thread 2
1 Byte	34.994756	33.090668
1 KB	5389.47368	1144.13408
1 MB	8206.542	1923.99266

	Random	
Block Size	Thread 1	Thread 2
1 Byte	6.4020487	6.4020487
1 KB	5120	6041.29794
1 MB	12000	19065.0182





- Analysis –

- Above figure represents the value of throughput for sequential and random read operation for 1 thread and 2 threads.
- Here nature of graphs is different like in first graph the throughput for 1 thread is high than 2 threads, whereas in second graph throughput of 1 thread is lower than 2 threads.

- Write Operations (Throughput)

	Thread 1	
Block Size	Sequential	Random
1 Byte	32.938076	0.846762
1 KB	1163.63636	572.067039
1 MB	1978.44528	1978.44528

	Thread 2	
Block Size	Sequential	Random
1 Byte	32.615786	0.845802
1 KB	1125.27473	567.313019
1 MB	1807.88966	1823.61044


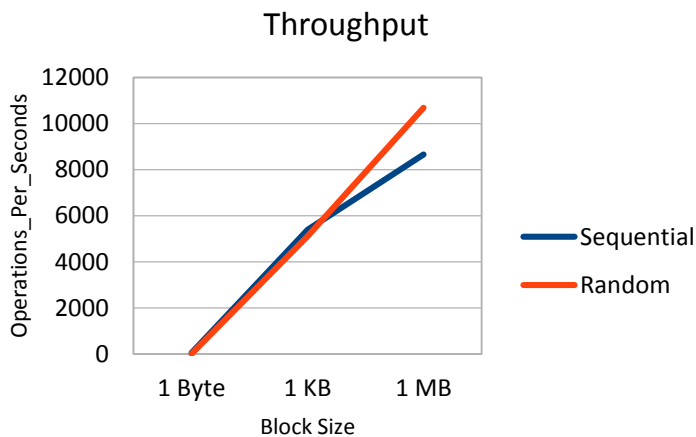


- Analysis-
  - Table and Graph above represent Throughput for Sequential and Random disk write operations
  - From above graph we can conclude that throughput keeps increasing with increase in block size, even it slightly varies for sequential and random operations. Its tread is same for sequential and Random.

• Read Operations (Throughput)

	Thread 1	
Block Size	Sequential	Random
1 Byte	44.99476	6.4020487
1 KB	5389.47368	5120.554
1 MB	8651.2545	10689.65

	Thread 2	
Block Size	Sequential	Random
1 Byte	33.090668	0.825631
1 KB	1144.13408	604.129794
1 MB	1923.99266	1906.50182





- Analysis-

- Table and Graph above represent Throughput for Sequential and Random disk read operations
- From above graph we can conclude that throughput keeps increasing with increase in block size, even it slightly varies for sequential and random operations. Its trend is same for sequential and Random.

- Write Operations (Latency)

	Thread 1	
	Sequential	Random
1 Byte	30.36	1180.97
1 KB	0.859375	1.748047
1 MB	0.505447	0.505447

	Thread 2	
	Sequential	Random
1 Byte	30.66	1182.31
1 KB	0.888672	1.762695
1 MB	0.553131	0.548363





- Analysis-

- Table and Graph above represent Latency for 1 Thread and 2 threads for disk write operations
- From above graph we can conclude that Latency is high when we have block size of 1 Byte and it gets lower with increase in size of block. (In case of Random write operations)


- Read Operations (Latency)

	Thread 1	
	Sequential	Random
1 Byte	200.21	15.62
1 KB	0.185547	0.019531
1 MB	0.11053	0.019093

	Thread 2	
	Sequential	Random
1 Byte	30.22	1211.195
1 KB	0.874023	1.655273
1 MB	0.519753	0.524521



- Analysis-
  - Table and Graph above represent Latency for 1 Thread and 2 threads for disk read operations
  - From above graph we can conclude that Latency is high when we have block size of 1 Byte and it gets lower with increase in size of block.
- Disk Theoretical Performance –**
  - As mentioned on Amazon Web Service Site Disk has performance of 160 mbps
- Extra Credit**
  - Here we have compared our results with stream benchmark. Below is screen shot for stream systems outputs.
  - So here we have achieved almost below % of ideal value compared with stream
  - Read + Write Operation = 115%



```
ubuntu@ip-172-31-56-58: ~  
ubuntu@ip-172-31-56-58:~$ ls  
stream.c  
ubuntu@ip-172-31-56-58:~$ sudo chmod 400 stream.c  
ubuntu@ip-172-31-56-58:~$ gcc -o streamapp stream.c  
ubuntu@ip-172-31-56-58:~$ ./streamapp  
-----  
STREAM version $Revision: 5.10 $  
-----  
This system uses 8 bytes per array element.  
-----  
Array size = 10000000 (elements), Offset = 0 (elements)  
Memory per array = 76.3 MiB (= 0.1 GiB).  
Total memory required = 228.9 MiB (= 0.2 GiB).  
Each kernel will be executed 10 times.  
The *best* time for each kernel (excluding the first iteration)  
will be used to compute the reported bandwidth.  
-----  
Your clock granularity/precision appears to be 1 microseconds.  
Each test below will take on the order of 26612 microseconds.  
 (= 26612 clock ticks)  
Increase the size of the arrays if this shows that  
you are not getting at least 20 clock ticks per test.  
-----  
WARNING -- The above is only a rough guideline.  
For best results, please be sure you know the  
precision of your system timer.  
-----  
Function      Best Rate MB/s  Avg time     Min time     Max time  
Copy:         6034.6    0.026914    0.026514    0.027427  
Scale:        5955.9    0.027616    0.026864    0.028322  
Add:          8750.5    0.028082    0.027427    0.028816  
Triad:        8024.9    0.030805    0.029907    0.031178  
-----  
Solution Validates: avg error less than 1.000000e-13 on all three arrays  
-----  
ubuntu@ip-172-31-56-58:~$
```